

Discussion Notes:

Questions and Extensions from the AD(2024) Robustness Presentation

Huang Zheng

May 2026

These notes supplement the presentation (`results/beamer_presentation.pdf`) with mathematical derivations, additional experiments, and responses to questions from classmates and faculty.

Q1: Does shuffling concept labels actually change the Ridge fit?

The question: If you randomly relabel all 296 concepts within each meeting, the within-meeting distribution of sentiment scores is preserved. Does Ridge regression actually produce a different fit? Or is it just computing the same linear combination?

Short answer: For the original 296 concepts (86% pro-cyclical), the Ridge fit is nearly unchanged because the dominant singular vector of the design matrix captures a meeting-level “document tone” that is invariant to column permutation.

Mathematical Derivation

Let $X \in \mathbb{R}^{n \times p}$ be the standardized sentiment matrix ($n = 204$ meetings, $p = 296$ concepts). The key empirical fact is that the first singular value dominates:

$$X \approx \sigma_1 \mathbf{u}_1 \mathbf{v}_1^\top, \quad \frac{\sigma_1^2}{\sum_{i=1}^n \sigma_i^2} \approx 10.6\% \quad (1)$$

where $\mathbf{u}_1 \in \mathbb{R}^n$ captures meeting-level document tone and $\mathbf{v}_1 \in \mathbb{R}^p$ captures concept-level loadings. The remaining $p - 1$ singular directions collectively explain 89.4% of the Frobenius norm but are individually small.

A column permutation $P \in \mathbb{R}^{p \times p}$ transforms the design matrix to $\tilde{X} = XP$. The Ridge estimator with penalty λ is:

$$\hat{\beta}(\lambda) = (X^\top X + \lambda I)^{-1} X^\top \mathbf{y} \quad (2)$$

Under the rank-1 approximation $X \approx \sigma \mathbf{u} \mathbf{v}^\top$:

$$X^\top X \approx \sigma^2 \mathbf{v} \mathbf{v}^\top \quad (3)$$

$$\tilde{X}^\top \tilde{X} = P^\top X^\top X P \approx \sigma^2 (P^\top \mathbf{v})(P^\top \mathbf{v})^\top \quad (4)$$

Both have rank 1 with the same dominant eigenvalue $\sigma^2 \|\mathbf{v}\|^2$. The Ridge fitted values are:

$$\hat{\mathbf{y}} = X(X^\top X + \lambda I)^{-1} X^\top \mathbf{y} \approx \frac{\sigma^2 \|\mathbf{v}\|^2}{\sigma^2 \|\mathbf{v}\|^2 + \lambda} \cdot \mathbf{u} \mathbf{u}^\top \mathbf{y} \quad (5)$$

This expression is **independent of $\mathbf{v}$** (the concept loadings). Permuting the columns ($\mathbf{v} \rightarrow P^\top \mathbf{v}$) does not change $\hat{\mathbf{y}}$. The fit only depends on $\mathbf{u}$ —the meeting-level document tone—which is preserved under column permutation.

What this means: Ridge is not learning 296 independent signals. Under strong regularization ($\lambda = 924$, selected by CV), it is essentially computing a **weighted average** of 296 noisy measurements of the same underlying document tone. Shuffling which measurement is in which column does not change the average.

Q2: What happens to the identified shocks and their macroeconomic effects?

The question: If we use the shuffled-label shocks in a Local Projections IRF, do they still produce the expected macroeconomic response? Is there a mathematical relationship?

Mathematical Derivation

Let $\varepsilon = \mathbf{y} - \hat{\mathbf{y}}(X)$ be the original shock and $\tilde{\varepsilon} = \mathbf{y} - \hat{\mathbf{y}}(\tilde{X})$ be the shock from a single shuffle. Write:

$$\tilde{\varepsilon} = \varepsilon + \boldsymbol{\eta}, \quad \text{where } \boldsymbol{\eta} = \hat{\mathbf{y}}(X) - \hat{\mathbf{y}}(\tilde{X}) \quad (6)$$

Key property: By construction of Ridge, $\varepsilon \perp \text{col}(X)$. Since $\boldsymbol{\eta} \in \text{col}(X) \cup \text{col}(\tilde{X})$, we have $\text{cov}(\varepsilon, \boldsymbol{\eta}) = 0$. Therefore:

$$\text{corr}(\varepsilon, \tilde{\varepsilon}) = \frac{\sigma_\varepsilon}{\sigma_{\tilde{\varepsilon}}} = \sqrt{\frac{\text{var}(\varepsilon)}{\text{var}(\tilde{\varepsilon})}} \quad (7)$$

Mean shuffle (100 runs): Let $\tilde{\varepsilon}^{(k)}$ be the shock from the k -th random permutation. The mean shuffled shock is:

$$\bar{\varepsilon} = \frac{1}{100} \sum_{k=1}^{100} \tilde{\varepsilon}^{(k)} = \varepsilon + \frac{1}{100} \sum_{k=1}^{100} \boldsymbol{\eta}^{(k)} \quad (8)$$

Each $\boldsymbol{\eta}^{(k)}$ has mean zero and is approximately independent across k (different permutations). By the law of large numbers, $\frac{1}{100} \sum \boldsymbol{\eta}^{(k)} \rightarrow 0$, so $\bar{\varepsilon} \approx \varepsilon$.

Empirical result: $\text{corr}(\varepsilon, \bar{\varepsilon}) = 0.96$, variance ratio ≈ 1.0 .

Implications for IRF

Consider a Local Projection:

$$y_{t+h} - y_{t-1} = \alpha_h + \beta_h \varepsilon_t + \boldsymbol{\gamma}^\top \mathbf{w}_t + e_{t+h} \quad (9)$$

With the shuffled shock $\tilde{\varepsilon}_t = \varepsilon_t + \eta_t$:

$$\text{plim } \tilde{\beta}_h = \frac{\text{cov}(\Delta y, \tilde{\varepsilon})}{\text{var}(\tilde{\varepsilon})} = \frac{\text{cov}(\Delta y, \varepsilon) + \text{cov}(\Delta y, \eta)}{\text{var}(\tilde{\varepsilon})} \quad (10)$$

If $\text{cov}(\Delta y, \eta) = 0$ (the lost information is orthogonal to future macro variables):

$$\tilde{\beta}_h = \frac{\text{var}(\varepsilon)}{\text{var}(\tilde{\varepsilon})} \cdot \beta_h \quad (11)$$

The IRF **shape** is identical; only the amplitude is scaled. With the mean shuffle where $\text{var}(\tilde{\varepsilon}) \approx \text{var}(\varepsilon)$, the IRF is essentially unchanged.

Empirical IRF Results

We estimate Local Projections (Jordà 2005) for the cumulative unemployment response at horizons $h = 0, 1, \dots, 48$ months:

$$u_{t+h} - u_{t-1} = \alpha_h + \beta_h \cdot \text{shock}_t + \sum_{\ell=1}^4 \gamma_{\ell} u_{t-\ell} + \sum_{\ell=1}^4 \delta_{\ell} \text{shock}_{t-\ell} + e_{t+h} \quad (12)$$

We compare three shock series: (1) the original AD(2024) shock, (2) the mean shuffled shock (100 random permutations), and (3) a forecasts-only shock (no sentiment indicators, $R^2 = 0.49$).




Figure 1: Cumulative unemployment response: original vs. shuffled shocks. The two IRFs are nearly identical in both shape and magnitude.

Local Projections: Unemployment Response to Monetary Policy Shocks




Figure 2: Four-panel IRF comparison. (A) Unemployment response to all three shocks. (B) Log unemployment. (C) Original vs. shuffled close-up: nearly identical. (D) Forecasts-only shock produces a smaller, noisier response—the shock is “dirtier” because it mixes systematic policy with true shocks.




Figure 3: Original vs. mean shuffled shock series. Correlation = 0.96.

Key findings:

- The original and shuffled shocks produce nearly identical IRFs—the peak response occurs at the same horizon (47 months) and the shapes are indistinguishable within confidence

bands.

- The forecasts-only shock produces a much smaller IRF because it is contaminated with the systematic component of monetary policy. This illustrates why AD(2024)'s sentiment-augmented approach is necessary for clean shock identification.
- The wide confidence bands at long horizons are a well-known feature of Local Projections and reflect the accumulation of forecast error.

—

Q3: What is the intrinsic dimension of the model? Should we try Random Projections?

The question: With 3,224 variables on 204 observations, is the model overfitting? How many independent dimensions does the data actually have?

Five Diagnostic Tests

Table 1: Overfitting Diagnostics for the Full Specification

Test	Method	Key Result	Conclusion
A. Random Projection	Project X to k random directions	$k = 1500$ (47%) needed for 90% R^2	Signal not in very low dimension, but far below $p = 3,224$
B. Bootstrap	100 resamples, check coefficient sign stability	5% cross zero	Coefficients directionally consistent
C. Noise Injection	Add k columns of pure noise	+10% noise $\Delta R^2 = +0.002$	Regularization prevents noise-fitting
D. Random Drop	Drop 50% of sentiment columns	$\Delta R^2 = +0.075$	Variables contain genuine information
E. Residual AC	Ljung-Box, Durbin-Watson	DW = 1.84	Residuals are white noise

Random Projection




Figure 4: Random Projection: projecting the 3,224-dimensional design matrix onto k random Gaussian directions. The baseline $R^2 = 0.96$ (red dashed line). Reaching 90% of baseline R^2 requires $k = 1500$ dimensions (47% of the original p).

Effective Rank

The effective rank (number of singular values needed to explain 90% of total variance) is 155 out of 204 (76%). This tells us:

- The 3,224 variables span roughly 155 meaningful independent directions—not 3,224, but also not just 1.
- Ridge with $\lambda = 924$ shrinks the remaining ~ 50 directions to near zero.
- This is **not** overfitting: the residuals are white noise, noise injection does not inflate R^2 , and bootstrap coefficients are stable.

Bottom line: The model is highly redundant (3,224 variables $\rightarrow$ ~ 155 effective dimensions), but Ridge regularization successfully handles this redundancy. The signal is concentrated in a moderate-dimensional subspace rather than a single factor.

Q4: What if we used a different concept list?

The question: The concept list is 86% pro-cyclical. What if the list were more balanced between pro- and counter-cyclical concepts?

Experiment: Same Size, Varying Composition

We construct three concept lists of equal size (68 concepts each), all scored by the same LM dictionary within ± 10 -word windows. The only difference is the pro-/counter-cyclical composition:

Table 2: Same Size, Different Composition

List (68 concepts each)	Pro-/Counter-	Full R^2	CV α
100% Pro-cyclical	68p / 0c	0.86	574
Random (mirrors original)	57p / 11c	0.83	924
50/50 Balanced	34p / 34c	0.75	1,487
<i>For reference:</i>			
Original 296	256p / 34c	0.96	924
Balanced 296 (LLM)	148p / 148c	1.00	≈ 1

Why the CV α Changes

The Ridge CV objective can be written in the SVD basis as:

$$\text{CV}(\alpha) \propto \sum_{i=1}^n \frac{\sigma_i^2}{(\sigma_i^2 + \alpha)^2} (\mathbf{u}_i^\top \mathbf{y})^2 \quad (13)$$

When $\sigma_1 \gg \sigma_2$ (strong common factor, as in 100% pro-cyclical), the optimal α is large because shrinking higher dimensions costs little in fit but greatly reduces variance. When the singular values are more evenly distributed (50/50 balanced), each dimension appears to carry independent signal, so CV selects a small α .

For the balanced 296-concept list, the singular value spectrum is flat enough that CV selects $\alpha \approx 1$ —essentially no regularization. With $p/n \approx 15.8$, this produces $R^2 = 1.00$, a clear case of overfitting.

Key insight: AD(2024)’s method is robust *conditional on the concept list being dominated by a common factor*. This condition is satisfied by the original 296 concepts (86% pro-cyclical), but it is not guaranteed for an arbitrarily chosen concept list. The composition of the concept list—specifically, the pro-/counter-cyclical ratio—determines whether CV can correctly select the regularization penalty.

Q5: LLM 595 vs. Original 296 Benchmark

For completeness, we benchmark the LLM-identified 595 concepts against the original 296, using identical dictionary scoring:

Table 3: Benchmark: Original 296 vs. LLM 595 (Mid Specification)

Specification	Variables	R^2
Forecasts only (no sentiment)	132	0.49
+ Original 296 sentiments	428	0.64
+ LLM 595 sentiments	727	0.68

The LLM list slightly outperforms the original in the mid specification (+0.05 R^2). However, the full specification (with lags and nonlinear terms) is not comparable: the LLM’s 595 concepts produce 6,214 variables, causing $R^2 = 1.00$ (overfitting). This is consistent with the finding above: the LLM list is more balanced (52/42), leading to a flatter singular value spectrum and weaker regularization.

Summary of Key Results

1. **Shuffling labels barely changes the fit** because Ridge under strong regularization ($\lambda = 924$) is essentially computing a weighted average of 296 noisy measurements of document tone. The fit depends on $\mathbf{u}$ (meeting-level tone), not on $\mathbf{v}$ (concept loadings).
2. **Mean shuffled shocks are nearly identical to original shocks** ($r = 0.96$). The IRF produced by the two shock series are essentially the same. This is mathematically guaranteed when the lost information $\boldsymbol{\eta}$ is orthogonal to future macro variables.
3. **The model does not overfit** despite $p/n = 15.8$. Ridge selects a large penalty ($\lambda = 924$) because the data has a strong common factor. Five diagnostic tests confirm this.
4. **List composition determines regularization.** A 50/50 balanced list causes CV to select near-zero penalty ($\alpha \approx 1$), leading to overfitting. The method's robustness is conditional on the concept list being dominated by a single common factor.
5. **Concept selection matters at the level of list composition, not individual words.** Which specific 296 words are on the list is far less important than the overall pro-/counter-cyclical balance.

All code is in `src/`. Results and figures are in `results/`. The presentation is at `results/beamer_presentation.pdf`.